

Statistical Learning

Lecture 07b - Some Non-Linearity

ANU - RSFAS

Last Updated: Thu Apr 21 13:02:44 2022

Moving Beyond Linearity

- The truth is never linear! Or almost never!
- But often the linearity assumption is good enough.
- When its not . . .
 - polynomials,
 - step functions,
 - splines,
 - local regression, and
 - generalized additive models

offer a lot of flexibility, without losing the ease and interpretability of linear models.

Polynomial Regression

- Create new variables $X_1 = X$, $X_2 = X^2$, etc and then treat as multiple linear regression.
- May not be interested in the coefficients; but more interested in the fitted function values at any value x_0 :

$$\hat{f}(x_0) = \hat{\beta}_0 + \hat{\beta}_1 x_0 + \hat{\beta}_2 x_0^2 + \hat{\beta}_3 x_0^3 + \cdots \hat{\beta}_k x_0^k$$

- Since $\hat{f}(x_0)$ is a linear function of the $\hat{\beta}$ s, we can get a simple expression for **pointwise-variances**

$$\text{Var}[\hat{f}(x_0)]$$

at any value x_0 .

- The we can compute **confidence intervals** (or we could use the bootstrap):

$$\hat{f}(x_0) \pm 1.96 \text{ se}[\hat{f}(x_0)]$$

- Recall the matrix form!

$$\mathbf{x}_0^{*'} = (1, x_0, x_0^2, \dots, x_0^k)$$

$$\hat{\beta} = (\hat{\beta}_0, \hat{\beta}_1, \hat{\beta}_2, \dots, \hat{\beta}_k)$$

$$\begin{aligned} \text{Var}[\hat{f}(x_0)] &= \text{Var}[\mathbf{x}_0^{*'} \hat{\beta}] = \mathbf{x}_0^{*'} V[\hat{\beta}] \mathbf{x}_0^* \\ &= \sigma^2 \mathbf{x}_0^{*'} (\mathbf{X}' \mathbf{X})^{-1} \mathbf{x}_0^* \\ &\Rightarrow \hat{\sigma}^2 \mathbf{x}_0^{*'} (\mathbf{X}' \mathbf{X})^{-1} \mathbf{x}_0^* \end{aligned}$$

$$\text{se}[\hat{f}(x_0)] = \sqrt{\text{Var}[\hat{f}(x_0)]}$$

- Let's consider the following data set on Wages (3000 workers).

```
library(ISLR)
data(Wage)
attach(Wage)
head(Wage)
```

##	year	age	maritl	race	education
## 231655	2006	18	1. Never Married	1. White	1. < HS Grad
## 86582	2004	24	1. Never Married	1. White	4. College Grad
## 161300	2003	45	2. Married	1. White	3. Some College
## 155159	2003	43	2. Married	3. Asian	4. College Grad
## 11443	2005	50	4. Divorced	1. White	2. HS Grad
## 376662	2008	54	2. Married	1. White	4. College Grad


##	region	jobclass	health
## 231655	2. Middle Atlantic	1. Industrial	1. <=Good
## 86582	2. Middle Atlantic	2. Information	2. >=Very Good
## 161300	2. Middle Atlantic	1. Industrial	1. <=Good
## 155159	2. Middle Atlantic	2. Information	2. >=Very Good
## 11443	2. Middle Atlantic	2. Information	1. <=Good
## 376662	2. Middle Atlantic	2. Information	2. >=Very Good

##	health_ins	logwage	wage
## 231655	2. No	4.318063	75.04315
## 86582	2. No	4.255273	70.47602
## 161300	1. Yes	4.875061	130.98218
## 155159	1. Yes	5.041393	154.68529
## 11443	1. Yes	4.318063	75.04315
## 376662	1. Yes	4.845098	127.11574

```
mod <- lm(wage ~ poly(age, degree = 4))

plot(age, wage, pch = 16, col = "gray",
      cex.lab = 2)
x <- seq(18, 80)

pred.y <- predict(mod, data.frame(age = x),
                  se.fit = TRUE)
lines(x, pred.y$fit, col = "blue", lwd = 2)
lines(x, pred.y$fit + 1.96 * pred.y$se.fit,
      col = "blue", lwd = 2, lty = 2)
lines(x, pred.y$fit - 1.96 * pred.y$se.fit,
      col = "blue", lwd = 2, lty = 2)
```



Logistic Regression

- Consider logistic regression. For example, we model

$$Pr(y_i > 250|x_i) = \frac{\exp(\hat{\beta}_0 + \hat{\beta}_1 x_0 + \hat{\beta}_2 x_0^2 + \hat{\beta}_3 x_0^3 + \cdots \hat{\beta}_k x_0^k)}{1 + \exp(\hat{\beta}_0 + \hat{\beta}_1 x_0 + \hat{\beta}_2 x_0^2 + \hat{\beta}_3 x_0^3 + \cdots \hat{\beta}_k x_0^k)}$$

- To get confidence intervals, compute upper and lower bounds on the **logit scale**, and then invert to get on **probability scale**.
- Recall on the logit scale we have a linear function of the $\hat{\beta}$ s:

$$\log\left(\frac{\widehat{\pi_i}}{1 - \pi_i}\right) = \hat{\eta}_i = \hat{\beta}_0 + \hat{\beta}_1 x_0 + \hat{\beta}_2 x_0^2 + \hat{\beta}_3 x_0^3 + \cdots \hat{\beta}_k x_0^k$$

- Caveat: polynomials have notorious tail behavior — very bad for extrapolation.
- We can fit these models using $y \sim \text{poly}(x, \text{degree} = k)$ in an R formula.

```
mod <- glm(I(wage > 250) ~ poly(age,
  degree = 4), family = "binomial")

pred.y <- predict(mod, data.frame(age = x),
  se.fit = TRUE)


pfit <- exp(pred.y$fit)/(1 + exp(pred.y$fit))

se.bands.logit <- cbind(pred.y$fit +
  1.96 * pred.y$se.fit, pred.y$fit -
  1.96 * pred.y$se.fit)



se.bands <- exp(se.bands.logit)/(1 +
  exp(se.bands.logit))

plot(age, I(wage > 250), type = "n",
  ylim = c(0, 0.2), cex.lab = 1.5)

lines(x, pfit, col = "blue", lwd = 2)
matlines(x, se.bands, col = "blue",
  lty = 2, lwd = 2)
```



Degree-4 Polynomial



Step Functions

- Another way of creating transformations of a variable — cut the variable into distinct regions.

$$C_1(X) = \mathbb{I}(X < 35), C_2(X) = \mathbb{I}(35 \leq X < 50), \dots, C_k(X) = \mathbb{I}(X > 65)$$



- Easy to work with. Creates a series of dummy variables representing each group.
- Useful way of creating interactions that are easy to interpret. For example, interaction effect of **Year** and **Age**:

$$\mathbb{I}(\text{Year} < 2005) \times \text{Age}, \quad \mathbb{I}(\text{Year} \geq 2005) \times \text{Age}$$

would allow for different linear functions in each age category.

- In R: `I(year<2005)` or `cut(age,c(18,25,40,65,90))`.
- The break points are called **knots**.
- Choice of cut-points or knots can be problematic. For creating nonlinearities, smoother alternatives such as **splines** are available.

Piecewise Constant



```
table(cut(age, 4))
```

```
##
```

```
## (17.9,33.5]    (33.5,49]    (49,64.5] (64.5,80.1]
```

```
##           750           1399           779           72
```

```
mod <- lm(wage ~ cut(age, 4))
```

```
plot(age, wage, pch = 16, col = "gray",  
      cex.lab = 2)
```


```
x <- seq(18, 80)
```

```
pred.y <- predict(mod, data.frame(age = x),  
                  se.fit = TRUE)
```

```
lines(x, pred.y$fit, col = "blue", lwd = 2)
```

```
lines(x, pred.y$fit + 1.96 * pred.y$se.fit,  
      col = "blue", lwd = 2, lty = 2)
```

```
lines(x, pred.y$fit - 1.96 * pred.y$se.fit,  
      col = "blue", lwd = 2, lty = 2)
```


Let's combine the Ideas - Piece-wise Polynomials

- Instead of a single polynomial in X over its whole domain, we can rather use different polynomials in regions defined by knots.


$$y_i = \begin{cases} \beta_{01} + \beta_{11}x_i + \beta_{21}x_i^2 + \beta_{31}x_i^3 + \epsilon_i & \text{if } x_i < c \\ \beta_{02} + \beta_{12}x_i + \beta_{22}x_i^2 + \beta_{32}x_i^3 + \epsilon_i & \text{if } x_i \geq c \end{cases}$$

- Better to add constraints to the polynomials, e.g. continuity.
- **Splines** have the “maximum” amount of continuity.


Piecewise Cubic




Continuous Piecewise Cubic



Cubic Spline



Linear Spline



Linear Splines

- A linear spline with knots at ξ_k , $k = 1, \dots, K$ is a piece-wise linear polynomial continuous at each knot.
- We can represent this model as



$$y_i = \beta_0 + \beta_1 b_1(x_i) + \beta_2 b_2(x_i) + \dots + \beta_{K+1} b_{K+1}(x_i) + \epsilon_i$$

where the b_k are **basis functions**.


$$\begin{aligned} b_1(x_i) &= x_i \\ b_{k+1}(x_i) &= (x_i - \xi_k)_+, \quad k = 1, \dots, K \end{aligned}$$

- Here, the $()_+$ means the positive part:


$$(x_i - \xi_k)_+ = \begin{cases} x_i - \xi_k & \text{if } x_i > \xi_k \\ 0 & \text{otherwise} \end{cases}$$




Piecewise Constant




Piecewise Linear



Continuous Piecewise Linear



Piecewise-linear Basis Function



- But we like something that is continuous (bottom left panel).
 - Let's add a constraint!
 - We want the function to the left of ξ_1 to be equal to the function to the right of ξ_1 .

$$f(\xi_1^-) = f(\xi_1^+)$$

- A direct way to do this, is to use a basis function which satisfies this constraint.

$$b_0(x_i) = 1, \quad b_1(x_i) = x_i, \quad b_2(x_i) = (x_i - \xi_1)_+, \quad b_3(x_i) = (x_i - \xi_2)_+$$

Cubic Spline

- With knots at ξ_k , $k = 1, \dots, K$ is a piece-wise cubic polynomial with continuous derivatives up to order 2 at each knot.



$$y_i = \beta_0 + \beta_1 b_1(x_i) + \beta_2 b_2(x_i) + \dots + \beta_{K+3} b_{K+3}(x_i) + \epsilon_i$$

- Again we can represent this model with truncated power basis functions:

$$\begin{aligned} b_1(x_i) &= x_i \\ b_2(x_i) &= x_i^2 \\ b_3(x_i) &= x_i^3 \\ b_{k+3}(x_i) &= (x_i - \xi_k)_+^3, \quad k = 1, \dots, K \end{aligned}$$


- Where:

$$(x_i - \xi_k)_+^3 = \begin{cases} (x_i - \xi_k)^3 & \text{if } x_i > \xi_k \\ 0 & \text{otherwise} \end{cases}$$




Piecewise Cubic Polynomials


Discontinuous




Continuous



Continuous First Derivative



Continuous Second Derivative

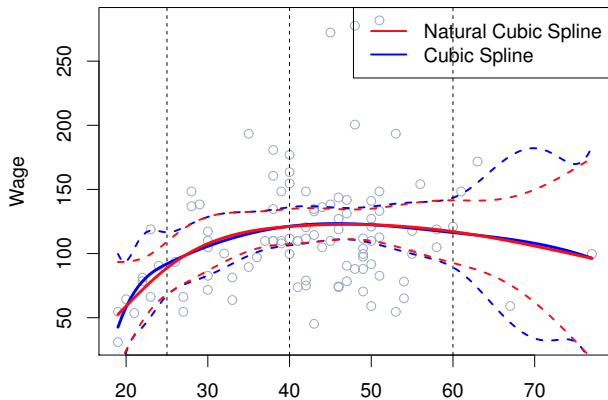


- The bottom right panel is represented by:

$$b_0(x_i) = 1, \quad b_1(x_i) = x_i, \quad b_2(x_i) = x_i^2, \quad b_3(x_i) = x_i^3$$
$$b_4(x_i) = (x_i - \xi_1)_+^3, \quad b_5(x_i) = (x_i - \xi_2)_+^3$$


Natural Cubic Splines

- A **natural cubic spline** extrapolates linearly beyond the boundary knots. This means there is an extra constraint!



- Fitting splines in R is easy: `bs(x, ...)` for any degree splines, and `ns(x, ...)` for natural cubic splines, in package `splines2`. Actually the package we want is `splines`, but only archive versions are available, but we can get `splines` by installing `splines2`.

Natural Cubic Spline



- The default is for a cubic spline so the **degree=3**.
- We can set the **knots**.
- Or we can set **degrees-of-freedom (df)** we want to use. We need to set them at least as large as the **degree**.
- The number of **dfs** over the degree is the number of **knots**.
- If **degree=3** and **df=4** then there is one knot at the median.
- If **degree=3** and **df=5** then there are two knots at the 1/3 and 2/3 quantiles.

`bs(x, df = NULL, knots = NULL, degree = 3, intercept = FALSE,
Boundary.knots = range(x))`


```
library(splines)

set.seed(1001)
x <- sort(rnorm(10))

test <- bs(x, degree = 3, df = 5)
attr(test, "knots")
```

```
## 33.33333% 66.66667%
```

```
## -0.6229437 -0.1775473
```


- Note: for `ns()` you cannot change the degree, it is set at 3.

```
mod1 <- lm(wage ~ bs(age, degree = 3,  
  df = 5))  
mod2 <- lm(wage ~ ns(age, df = 5))
```


```
plot(age, wage, pch = 16, col = "gray")
x <- seq(18, 80)

pred.y <- predict(mod1, data.frame(age = x),
  se.fit = TRUE)
lines(x, pred.y$fit, col = "blue")
lines(x, pred.y$fit + 1.96 * pred.y$se.fit,
  col = "blue")
lines(x, pred.y$fit - 1.96 * pred.y$se.fit,
  col = "blue")

pred.y2 <- predict(mod2, data.frame(age = x),
  se.fit = TRUE)
lines(x, pred.y2$fit, col = "red")
lines(x, pred.y2$fit + 1.96 * pred.y2$se.fit,
  col = "red")
lines(x, pred.y2$fit - 1.96 * pred.y2$se.fit,
  col = "red")
```



- Comparison of a degree-14 polynomial and a natural cubic spline, each with 15df.
- `ns(age, df=14)`
- `poly(age, deg=14)`



Smoothing Splines - More Fun!

- Consider this criterion for fitting a smooth function $g(x)$ to some data:

$$\underset{g \in \mathcal{S}}{\text{minimize}} \sum_{i=1}^n (y_i - g(x_i))^2 + \lambda \int g''(t)^2 dt$$

- The first term is the RSS, and tries to make $g(x)$ match the data at each x_i .
- The second term is a roughness penalty and controls how wiggly $g(x)$ is.
- It is modulated by the tuning parameter $\lambda \geq 0$.
 - The smaller λ , the more wiggly the function, eventually interpolating y_i when $\lambda = 0$.
 - As $\lambda \rightarrow \infty$ the function $g(x_i)$ becomes linear.

- The solution to the minimization is a natural cubic spline, with a knot at every unique value of x_i . The roughness penalty still controls the roughness via λ .
- Smoothing splines avoid the knot-selection issue, leaving a single λ to be chosen.
- The algorithmic details are too complex to describe here. In R, the function `smooth.spline()` will fit a smoothing spline.
- The vector of n fitted values can be written (\mathbf{S} is just the hat matrix \mathbf{H}):

$$\hat{\mathbf{g}}_\lambda = \mathbf{S}_\lambda \mathbf{y}$$

- The effective degrees of freedom are given by (note that is just the trace of the matrix):

$$df_{\lambda} = \sum_{i=1}^n \{\mathbf{S}_{\lambda}\}_{ii}$$

- We can specify **df** rather than λ ! In R: `smooth.spline(age, wage, df = 10)`
- The leave-one-out (LOO) cross-validated error is given by

$$CV_{LOOCV} = \sum_{i=1}^n \left[\frac{y_i - \hat{g}_{\lambda}^{(-i)}(x_i)}{1 - \{\mathbf{S}_{\lambda}\}_{ii}} \right]^2$$

- In R: `smooth.spline(age, wage, cv=TRUE)`

```
fit <- smooth.spline(age, wage, df = 16)
fit$lambda
```

```
## [1] 0.0006537868
```

```
fit2 <- smooth.spline(age, wage, cv = TRUE)
fit2$df
```


```
## [1] 6.794596
```

```
plot(age, wage, pch = 16, col = "gray")
```


```
lines(fit, col = "red ", lwd = 2)
```

```
lines(fit2, col = " blue", lwd = 2)
```

```
legend("topright", legend = c("16 DF",  
    "6.8 DF"), col = c("red", "blue"),  
    lty = 1, lwd = 2, cex = 0.8)
```



Local Regression - Yet Another Approach



- With a sliding weight function, we fit separate linear fits over the range of X by weighted least squares. Use the `loess()` function in R.

```
fit <- loess(wage ~ age, span = 0.2,  
            data = Wage)  
fit
```

```
## Call:  
## loess(formula = wage ~ age, data = Wage, span = 0.2)  
##  
## Number of Observations: 3000  
## Equivalent Number of Parameters: 16.42  
## Residual Standard Error: 39.92
```

```
fit2 <- loess(wage ~ age, span = 0.5,  
             data = Wage)  
fit2
```


```
## Call:  
## loess(formula = wage ~ age, data = Wage, span = 0.5)  
##  
## Number of Observations: 3000  
## Equivalent Number of Parameters: 7.13  
## Residual Standard Error: 39.89
```

```
plot(age, wage, pch = 16, col = "gray")
age.grid <- 20:80

pred1 <- predict(fit, data.frame(age = age.grid),
  se = TRUE)
lines(age.grid, pred1$fit, col = "red ",
  lwd = 2)
lines(age.grid, pred1$fit + 1.96 * pred1$se.fit,
  col = "red ", lwd = 2)
lines(age.grid, pred1$fit - 1.96 * pred1$se.fit,
  col = "red ", lwd = 2)

pred2 <- predict(fit2, data.frame(age = age.grid),
  se = TRUE)
lines(age.grid, pred2$fit, col = "blue ",
  lwd = 2)
lines(age.grid, pred2$fit + 1.96 * pred2$se.fit,
  col = "blue", lwd = 2)
lines(age.grid, pred2$fit - 1.96 * pred2$se.fit,
  col = "blue", lwd = 2)

legend("topright", legend = c("Span=0.2",
  "Span=0.5"), col = c("red", "blue"),
  lty = 1, lwd = 2)
```






Generalized Additive Models

- Allows for flexible nonlinearities in several variables, but retains the additive structure of linear models.

$$y_i = \beta_0 + f_1(x_{i1}) + f_2(x_{i2}) + \cdots + f_p(x_{ip}) + \epsilon_i$$

- You can fit a GAM using `gam()` — you can also just use `lm()` when using `bs()` or `ns()`.

```
library(gam)
mod <- gam(wage ~ ns(year, df = 5) +
           ns(age, df = 5) + education)
par(mfrow = c(1, 3))
plot(mod, se = TRUE, col = "red", cex.lab = 2)
```



- Can mix terms — some linear, some nonlinear — and use `anova()` to compare models.




```
mod <- gam(wage ~ ns(year, df = 5) +  
  ns(age, df = 5) + education)  
mod2 <- gam(wage ~ ns(age, df = 5) +  
  education)  
  
anova(mod2, mod)
```

```
## Analysis of Deviance Table  
##  
## Model 1: wage ~ ns(age, df = 5) + education  
## Model 2: wage ~ ns(year, df = 5) + ns(age, df = 5) + education  
##   Resid. Df Resid. Dev Df Deviance Pr(>Chi)  
## 1      2990      3712881  
## 2      2985      3690702  5      22179 0.003025 **  
## ---  
## Signif. codes:  
## 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

- Can use smoothing splines or local regression as well:

```
mod <- gam(wage ~ s(year, df = 5) +  
           lo(age, span = 0.5) + education)  
par(mfrow = c(1, 3))  
plot(mod, se = TRUE, col = "red", cex.lab = 2)
```

- Can use smoothing splines or local regression as well:






- GAMs are additive, although low-order interactions can be included in a natural way using, e.g. bivariate smoothers or interactions of the form `ns(age,df=5):ns(year,df=5)`.

GAMs for Classification

$$\log \left(\frac{Pr(x_i)}{1 - Pr(x_i)} \right) = \beta_0 + f_1(x_{i1}) + f_2(x_{i2}) + \cdots + f_p(x_{ip})$$

```
mod <- gam(I(wage > 250) ~ year + s(age,  
    df = 5) + education, family = binomial)  
par(mfrow = c(1, 3))  
plot(mod, se = TRUE, col = "red")
```

```
table(education, I(wage > 250))
```

```
##  
## education          FALSE TRUE  
## 1. < HS Grad        268    0  
## 2. HS Grad          966    5  
## 3. Some College     643    7  
## 4. College Grad     663   22  
## 5. Advanced Degree  381   45
```

- Let's fit the model again without the '1. < HS Grad' category.

```
mod <- gam(I(wage > 250) ~ year + s(age,  
  df = 5) + education, family = binomial,  
  subset = (education != "1. < HS Grad"))  
par(mfrow = c(1, 3))  
plot(mod, se = TRUE, col = "red")
```

